

lavaan: an R package for structural equation modeling and more

Version 0.3-1 (BETA)

Yves Rosseel
Department of Data Analysis
Ghent University (Belgium)

August 17, 2010

Abstract

The **lavaan** package is developed to provide useRs, researchers and teachers a free, open-source, but commercial-quality package for latent variable analysis. The long-term goal of **lavaan** is to implement all the state-of-the-art capabilities that are currently available in commercial packages, including support for various data types, discrete latent variables (aka mixture models) and multilevel datasets. Currently, the **lavaan** package provides support for confirmatory factor analysis, structural equation modeling, and latent growth curve models. In this document, we illustrate the use of **lavaan** by providing several examples. If you are new to **lavaan**, this is the first document to read.

1 Before you start

Before you start, read these points carefully:

- First of all, you must have a recent ($\geq 2.10.1$) version of **R** installed. You can download the latest version of **R** from this page: <http://cran.r-project.org/>.
- The **lavaan** package is not finished yet. But it is already very useful for most users, or so we hope. There are a number of known minor issues (see section 9), and some features are simply not implemented yet. Some important features that are currently *not* available in **lavaan** are:
 - support for categorical/censored variables (this will be available in the next release of **lavaan**)
 - support for discrete latent variables (mixture models)
 - support for hierarchical/multilevel datasets

We hope to add these features in the next year or so.

- We do not expect you to be an expert in **R**. In fact, the **lavaan** package was designed to be used by users that would normally never use **R**. Nevertheless, it may help to familiarize yourself a bit with **R**, just to be comfortable with it. Perhaps the most important skill that you may need to learn is how to import your own datasets (perhaps in an SPSS format) into **R**. There are many tutorials on the web to teach you just that. Once you have your data in **R**, you can start specifying your **lavaan** model. We have tried very hard to make it as easy as possible for users to fit their models. Of course, if you have suggestions on how we can improve things, please let us know.
- This document is written for first-time users (and beta-testers) of the **lavaan** package. It is not a reference manual, nor does it contain technical material on how things are done in the **lavaan** package. These documents are currently under preparation.
- The **lavaan** package is free open-source software. This means (among other things) that there is no warranty whatsoever.
- The numerical results of the **lavaan** package are typically very close, if not identical, to the results of the commercial package Mplus. If you wish to compare the results with other SEM packages, you can use the optional argument `mimic.Mplus=FALSE` when calling the `cfa`, `sem` or `growth` functions (see section 8.2).

2 Installation of the lavaan package

Since may 2010, the `lavaan` package is available on CRAN. Therefore, to install `lavaan`, simply start up R, and type:

```
> install.packages("lavaan")
```

You can check if the installation was succesful by typing

```
> library(lavaan)

This is lavaan 0.3-1
lavaan is BETA software! Please report any bugs.
```

If you see the startup message (showing the version number, and a reminder that this is beta software), you're all set. Move on to the next section. If you get an error, or nothing happens at all, please let us know. See section 11 for how to submit a bug report.

3 The model syntax

At the heart of the `lavaan` package is the 'model syntax'. The model syntax is a description of the model to be estimated. In this section, we briefly explain the elements of the `lavaan` model syntax. More details are given in the examples that follow.

In the R environment, a regression formula has the following form:

$$y \sim x1 + x2 + x3 + x4$$

In `lavaan`, a typical model is simply a set (or system) of regression formulas, where some variables (starting with an 'f' below) may be latent. For example:

$$\begin{aligned} y &\sim f1 + f2 + x1 + x2 \\ f1 &\sim f2 + f3 \\ f2 &\sim f3 + x1 + x2 \end{aligned}$$

If we have latent variables in any of the regression formulas, we need to 'define' them by listing their manifest indicators. We do this by using the special operator "`=~`", which can be read as *is manifested by*. For example, to define the three latent variabels `f1`, `f2` and `f3`, we can use something like:

$$\begin{aligned} f1 &=~ y1 + y2 + y3 \\ f2 &=~ y4 + y5 + y6 \\ f3 &=~ y7 + y8 + y9 + y10 \end{aligned}$$

Further more, variances and covariances are specified using a 'double tilde' operator, for example:

$$\begin{aligned} y1 &~~ y1 \\ y1 &~~ y2 \\ f1 &~~ f2 \end{aligned}$$

And finally, intercepts for observed and latent variables are simple regression formulas with only an intercept (explicitly denoted by the number '1') as the only predictor:

$$\begin{aligned} y1 &\sim 1 \\ f1 &\sim 1 \end{aligned}$$

Using these four *formula types*, a large variety of latent variable models can be described. But new formula types may be added in the near future. The current set of formula types is summarized in the table below.

formula type	operator	mnemonic
latent variable definition	<code>=~</code>	is measured by
regression	<code>~</code>	is regressed on
(residual) (co)variance	<code>~~</code>	is correlated with
intercept	<code>~ 1</code>	intercept

3.1 Entering the model syntax as a string literal

If the model syntax is fairly short, you can specify it interactively at the R prompt by enclosing the formulas with single quotes. For example:

```
> myModel <- ' # regressions
                y ~ f1 + f2 +
                  x1 + x2
                f1 ~ f2 + f3
                f2 ~ f3 + x1 + x2

                # latent variable definitions
                f1 =~ y1 + y2 + y3
                f2 =~ y4 + y5 + y6
                f3 =~ y7 + y8 +
                  y9 + y10

                # variances and covariances
                y1 ~~ y1
                y1 ~~ y2
                f1 ~~ f2

                # intercepts
                y1 ~ 1
                f1 ~ 1
                ,
```

This will produce a model syntax object, called `myModel` that can be used later when calling a function that actually estimates this model given a dataset. Note that formulas can be split over multiple lines, and you can use comments (starting with the `#` character) and blank lines within the single quotes to improve readability of the model syntax.

3.2 Reading the model syntax from an external file

If your model syntax is rather long, you may prefer to type it in a separate text file called, say, `myModel.lms`. This text file should be in a human readable format (not a Word document). Within R, you can then read the model syntax from the file as follows:

```
> myModel <- readLines("/mydirectory/myModel.lms")
```

The argument of `readLines` is the full path to the file containing the model syntax. Again, the model syntax object `myModel` can be used later to fit this model given a dataset.

4 Fitting latent variables models: two examples

4.1 A first example: confirmatory factor analysis (CFA)

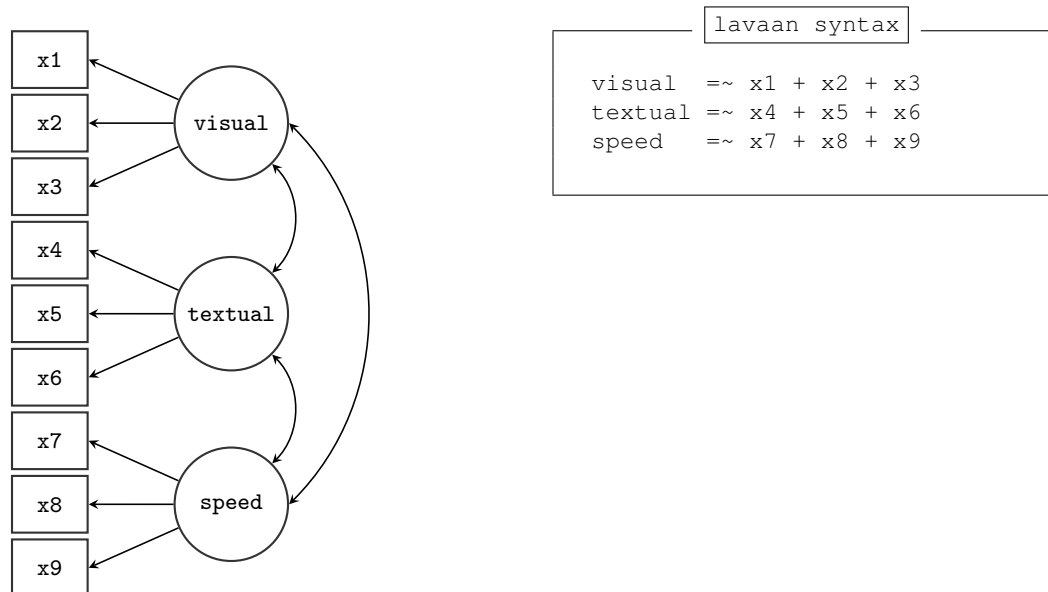
We start with a simple example of confirmatory factor analysis. The `lavaan` package contains a built-in dataset called `HolzingerSwineford1939`. See the help page for this dataset by typing

```
> ?HolzingerSwineford1939
```

at the R prompt. This is a ‘classic’ dataset that is used in many papers and books on Structural Equation Modeling (SEM), including some manuals of commercial SEM software packages. The data consists of mental ability test scores of seventh- and eighth-grade children from two different schools (Pasteur and Grant-White). In our version of the dataset, only 9 out of the original 26 tests are included. A CFA model that is often proposed for these 9 variables consists of three latent variables (or factors), each with three indicators:

- a *visual* factor measured by 3 variables: `x1`, `x2` and `x3`
- a *textual* factor measured by 3 variables: `x4`, `x5` and `x6`
- a *speed* factor measured by 3 variables: `x7`, `x8` and `x9`

The left panel of the figure below contains a simple graphical representation of the three-factor model. The right panel contains the corresponding `lavaan` syntax for specifying this model.



In this example, the model syntax only contains three ‘latent variable definitions’. Each formula has the following format:

```
latent variable =~ indicator1 + indicator2 + indicator3
```

We call these expressions *latent variable definitions* because they define how the latent variables are ‘manifested by’ a set of observed (or manifest) variables, often called ‘indicators’. Note that the special “=~” operator in the middle consists of a sign (“=”) character and a tilde (“~”) character next to each other. The reason why this model syntax is so short, is that behind the scenes, **lavaan** will take care of several things. First, by default, the factor loading of the first indicator of a latent variable is fixed to 1, thereby fixing the scale of the latent variable. Second, residuals variances are added automatically. And third, all latent variables are correlated by default. This way, the model syntax can be kept concise. On the other hand, the user remains in control, since all this ‘default’ behavior can be overridden. More on this later.

We can enter the model syntax using the single quotes:

```
> HS.model <- '
+   visual  =~ x1 + x2 + x3
+   textual =~ x4 + x5 + x6
+   speed   =~ x7 + x8 + x9
+ '
```

We can now fit the model as follows:

```
> fit <- cfa(HS.model, data = HolzingerSwineford1939)
```

The **lavaan** function **cfa** is a dedicated function for fitting confirmatory factor analysis models. The first argument is the user-specified model. The second argument is the dataset that contains the observed variables. Once the model has been fitted, the **summary** method provides a nice summary of the fitted model:

```
> summary(fit, fit.measures = TRUE)
```

Model converged normally after 35 iterations using ML

Minimum Function Chi-square	85.306
Degrees of freedom	24
P-value	0.0000

Chi-square test baseline model:

Minimum Function Chi-square	918.852
Degrees of freedom	36
P-value	0.0000

Full model versus baseline model:

Comparative Fit Index (CFI)	0.931
Tucker-Lewis Index (TLI)	0.896

Loglikelihood and Information Criteria:

Loglikelihood user model (H0)	-3737.745
Loglikelihood unrestricted model (H1)	-3695.092
Akaike (AIC)	7517.490
Bayesian (BIC)	7595.339

Root Mean Square Error of Approximation:

RMSEA	0.092
90 Percent Confidence Interval	0.071 0.114
P-value RMSEA <= 0.05	0.001

Standardized Root Mean Square Residual:

SRMR	0.065
------	-------

Model estimates:

	Estimate	Std.err	Z-value	P(> z)
Latent variables:				
visual =~				
x1	1.000			
x2	0.554	0.100	5.554	0.000
x3	0.729	0.109	6.685	0.000
textual =~				
x4	1.000			
x5	1.113	0.065	17.014	0.000
x6	0.926	0.055	16.703	0.000
speed =~				
x7	1.000			
x8	1.180	0.165	7.152	0.000
x9	1.082	0.151	7.155	0.000
Latent covariances:				
visual ~~				
textual	0.408	0.074	5.552	0.000
speed	0.262	0.056	4.660	0.000
textual ~~				
speed	0.173	0.049	3.518	0.000
Latent variances:				
visual	0.809	0.145	5.564	0.000
textual	0.979	0.112	8.737	0.000
speed	0.384	0.086	4.451	0.000
Residual variances:				
x1	0.549	0.114	4.833	0.000
x2	1.134	0.102	11.146	0.000
x3	0.844	0.091	9.317	0.000
x4	0.371	0.048	7.778	0.000
x5	0.446	0.058	7.642	0.000
x6	0.356	0.043	8.277	0.000
x7	0.799	0.081	9.823	0.000
x8	0.488	0.074	6.573	0.000
x9	0.566	0.071	8.003	0.000

The output should look familiar to users of other SEM software. If you find it confusing or esthetically unpleasing, again, please let us know, and we will try to improve it. To wrap up this first example, we summarize the code that was needed to fit this three-factor model:

```
R code

# load the lavaan package (only needed once per session)
library(lavaan)

# specify the model
HS.model <- ' visual  =~ x1 + x2 + x3
              textual =~ x4 + x5 + x6
              speed   =~ x7 + x8 + x9 '

# fit the model
fit <- cfa(HS.model, data=HolzingerSwineford1939)

# display summary output
summary(fit, fit.measures=TRUE)
```

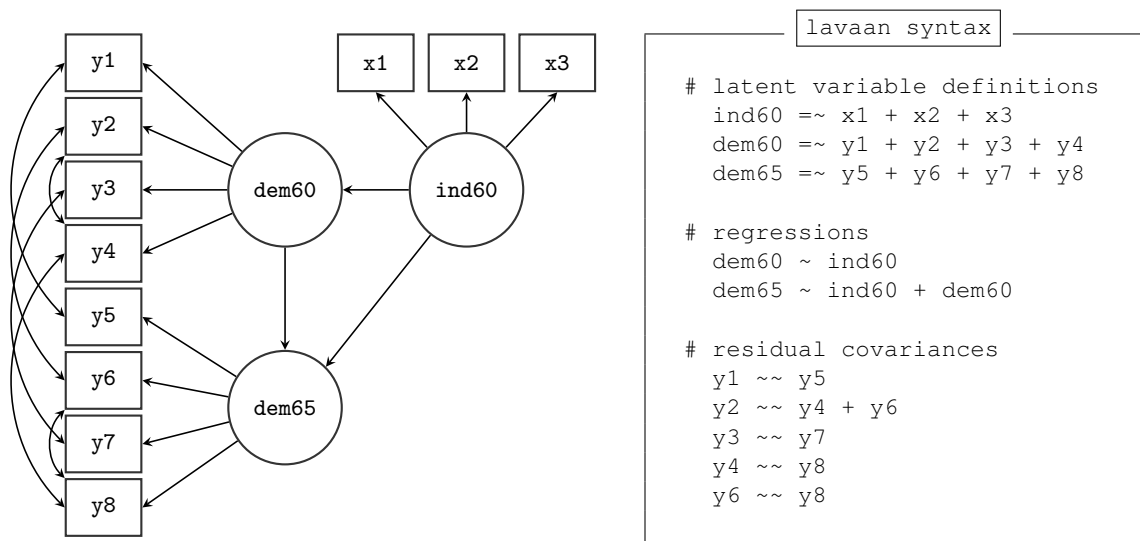
Simply copying this code and pasting it in **R** should work. The syntax illustrates the typical workflow in the **lavaan** package:

1. specify your model using the **lavaan** model syntax. In this example, only *latent variable definitions* have been used. In the following examples, other formula types will be used.
2. fit the model. This requires a dataset containing the observed variables (or alternatively the sample covariance matrix and the number of observations; see section 8.1). In this example, we have used the **cfa** function. Other functions in the **lavaan** package are **sem** and **growth** for fitting full structural equation models and growth curve models respectively.
3. extract information from the fitted model. This can be a long verbose summary, or it can be a single number only (say, the RMSEA value). In the spirit of **R**, you only get what you asked for. We do not print out unnecessary information that you would ignore anyway.

4.2 A second example: a structural equation model

In our second example, we will use the built-in **PoliticalDemocracy** dataset. This is a dataset that has been used by Bollen in his 1989 book on structural equation modeling (and elsewhere). To learn more about the dataset, see the help page and the references therein.

The left panel of the figure below contains a graphical representation of the model that we want to fit. The right panel contains the corresponding model syntax.



In this example, we use three different formula types: latent variable definitions, regression formulas, and (co)variance formulas. The regression formulas are similar to ordinary formulas in R. The (co)variance formulas typically have the following form:

variable ~~ variable

The variables can be either observed or latent variables. If the two variable names are the same, the expression refers to the variance (or residual variance) of that variable. If the two variable names are different, the expression refers to the (residual) covariance among these two variables. The *lavaan* package automatically makes the distinction between variances and residual variances.

In our example, the expression `y1 ~~ y5` allows the residual variances of the two observed variables to be correlated. This is sometimes done if it is believed that the two variables have something in common that is not captured by the latent variables. In this case, the two variables refer to identical scores, but measured in two different years (1960 and 1965 respectively). Note that the two expressions `y2 ~~ y4` and `y2 ~~ y6` can be combined into the expression `y2 ~~ y4 + y6`. This is just a shorthand notation.

We enter the model syntax as follows:

```
> model <- '
+   # measurement model
+   ind60 =~ x1 + x2 + x3
+   dem60 =~ y1 + y2 + y3 + y4
+   dem65 =~ y5 + y6 + y7 + y8
+
+   # regressions
+   dem60 ~ ind60
+   dem65 ~ ind60 + dem60
+
+   # residual correlations
+   y1 ~~ y5
+   y2 ~~ y4 + y6
+   y3 ~~ y7
+   y4 ~~ y8
+   y6 ~~ y8
+ '
```

To fit the model and see the results we can type:

```
> fit <- sem(model, data = PoliticalDemocracy)
> summary(fit, standardized = TRUE)
```

Model converged normally after 95 iterations using ML

Minimum Function Chi-square	38.125
Degrees of freedom	35
P-value	0.3292

	Estimate	Std.err	Z-value	P(> z)	Std.lv	Std.all
Latent variables:						
ind60 =~						
x1	1.000				0.670	0.920
x2	2.180	0.139	15.742	0.000	1.460	0.973
x3	1.819	0.152	11.967	0.000	1.218	0.872
dem60 =~						
y1	1.000				2.223	0.850
y2	1.257	0.182	6.888	0.000	2.794	0.717
y3	1.058	0.151	6.987	0.000	2.351	0.722
y4	1.265	0.145	8.722	0.000	2.812	0.846
dem65 =~						
y5	1.000				2.103	0.808
y6	1.186	0.169	7.024	0.000	2.493	0.746
y7	1.280	0.160	8.002	0.000	2.691	0.824
y8	1.266	0.158	8.007	0.000	2.662	0.828
Regressions:						
dem60 ~						
ind60	1.483	0.399	3.715	0.000	0.447	0.447

```

dem65 ~
  ind60          0.572    0.221    2.586    0.010    0.182    0.182
  dem60          0.837    0.098    8.514    0.000    0.885    0.885

Residual covariances:
  y1 ~~
  y5          0.624    0.358    1.741    0.082    0.624    0.092
  y2 ~~
  y4          1.313    0.702    1.870    0.061    1.313    0.101
  y6          2.153    0.734    2.934    0.003    2.153    0.165
  y3 ~~
  y7          0.795    0.608    1.308    0.191    0.795    0.075
  y4 ~~
  y8          0.348    0.442    0.787    0.431    0.348    0.033
  y6 ~~
  y8          1.356    0.568    2.386    0.017    1.356    0.126

Latent variances:
  ind60          0.448    0.087    5.173    0.000    1.000    1.000

Residual variances:
  x1          0.082    0.019    4.184    0.000    0.082    0.154
  x2          0.120    0.070    1.718    0.086    0.120    0.053
  x3          0.467    0.090    5.177    0.000    0.467    0.239
  y1          1.891    0.444    4.256    0.000    1.891    0.277
  y2          7.373    1.374    5.366    0.000    7.373    0.486
  y3          5.068    0.952    5.325    0.000    5.068    0.478
  y4          3.148    0.739    4.261    0.000    3.148    0.285
  y5          2.351    0.480    4.895    0.000    2.351    0.347
  y6          4.954    0.914    5.419    0.000    4.954    0.443
  y7          3.431    0.713    4.814    0.000    3.431    0.322
  y8          3.254    0.695    4.685    0.000    3.254    0.315
  dem60       3.956    0.921    4.294    0.000    0.800    0.800
  dem65       0.172    0.215    0.803    0.422    0.039    0.039

```

The function `sem` is very similar to the `cfa` function. In fact, the two functions are currently almost identical, but this may change in the future. In the `summary` method, we omitted the `fit.measures=TRUE` argument. Therefore, you only get the basic chi-square statistic. The argument `standardized=TRUE` augments the output with standardized parameter values. Two extra columns of standardized parameter values are printed. In the first column (labeled `Std.lv`), only the latent variables are standardized. In the second column (labeled `Std.all`), both latent and observed variables are standardized. The latter is often called the ‘completely standardized solution’.

The complete code to specify and fit this model is printed again below:


```

model <- '
# measurement model
ind60 =~ x1 + x2 + x3
dem60 =~ y1 + y2 + y3 + y4
dem65 =~ y5 + y6 + y7 + y8

# regressions
dem60 ~ ind60
dem65 ~ ind60 + dem60

# residual correlations
y1 ~~ y5
y2 ~~ y4 + y6
y3 ~~ y7
y4 ~~ y8
y6 ~~ y8
'

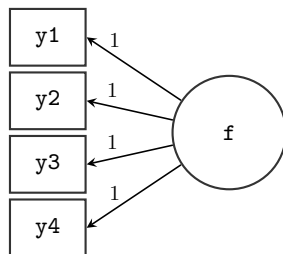
fit <- sem(model, data=PoliticalDemocracy)
summary(fit, standardized=TRUE)

```

5 Fixing parameters, starting values and equality constraints

5.1 Fixing parameters

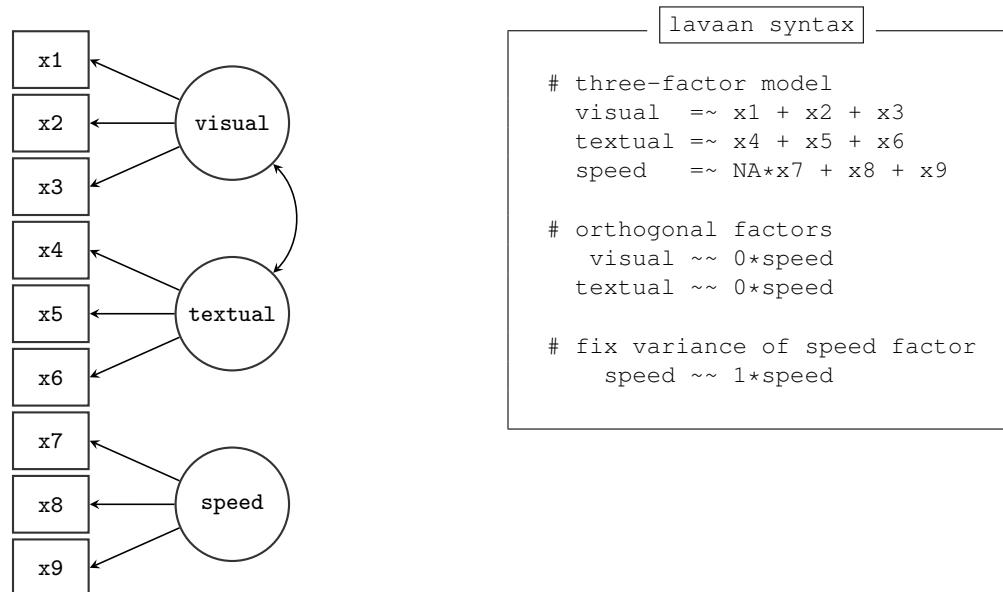
Consider a simple one-factor model with 4 indicators. By default, *lavaan* will always fix the factor loading of the first indicator to 1. The other three factor loadings are free, and their values are estimated by the model. But suppose that you have good reasons to fix all the factor loadings to 1. The syntax below illustrates how this can be done:



lavaan syntax

```
f =~ y1 + 1*y2 + 1*y3 + 1*y4
```

In general, to fix a parameter in a *lavaan* formula, you need to pre-multiply the corresponding variable in the formula by a numerical value. This is called the pre-multiplication mechanism and will be used for many purposes. As another example, consider again the three-factor Holzinger and Swineford CFA model. Recall that, by default, all latent variables in a CFA model are correlated. But if you wish to fix the correlation (or covariance) between a pair of latent variables to zero, you need to explicitly add a covariance-formula for this pair, and fix the parameter to zero. In the figure below, we allow the covariance between the latent variables **visual** and **textual** to be free, but the two other covariances are fixed to zero. In addition, we fix the variance of the **speed** factor to unity. Therefore, there is no need anymore to set the factor loading of its first indicator (**x7**) equal to one. To force this factor loading to be free, we pre-multiply it with **NA**, as a hint to *lavaan* that the value of this parameter is still unknown.



If you need to constrain all covariances of the latent variables in a CFA model to be orthogonal, there is a shortcut. You can omit the covariance formulas in the model syntax and simply add an `orthogonal=TRUE` argument to the `cfa` function call:

```
> HS.model <- ' visual =~ x1 + x2 + x3
+               textual =~ x4 + x5 + x6
+               speed  =~ x7 + x8 + x9 '
> fit.HS.ortho <- cfa(HS.model, data=HolzingerSwineford1939, orthogonal=TRUE)
```

Similarly, if you want to fix the variances of *all* the latent variables in a CFA model to unity, there is again a shortcut. Simply add a `std.lv=TRUE` argument to the `cfa` function call:

```
> HS.model <- ' visual =~ x1 + x2 + x3
+               textual =~ x4 + x5 + x6
+               speed  =~ x7 + x8 + x9 '
> fit <- cfa(HS.model, data=HolzingerSwineford1939, std.lv=TRUE)
```

If the `std.lv=TRUE` argument is used, the factor loadings of the first indicator of each latent variable will no longer be fixed to 1.

5.2 Starting values

The `lavaan` package automatically generates starting values for all free parameters. Normally, this works fine. But if you must provide your own starting values, you are free to do so. The way it works is based on the pre-multiplication mechanism that we discussed before. But the numeric constant is now the argument of a special function `start()`. An example will make this clear:

lavaan syntax

```
visual =~ x1 + start(0.8)*x2 + start(1.2)*x3
textual =~ x4 + start(0.5)*x5 + start(1.0)*x6
speed  =~ x7 + start(0.7)*x8 + start(1.8)*x9
```

The factor loadings of the first indicators (`x1`, `x4` and `x7`) are fixed, so no starting values are needed. But for all other factor loadings, starting values are provided in this example.

5.3 Parameter names

A nice property of the `lavaan` package is that all free parameters are automatically named according to a simple set of rules. This is convenient, for example, if equality constraints are needed (see the next subsection). To see how the naming mechanism works, we will use the model that we used for the Political Democracy data.

```

> model <- '
+   # latent variable definitions
+   ind60 =~ x1 + x2 + x3
+   dem60 =~ y1 + y2 + y3 + y4
+   dem65 =~ y5 + y6 + y7 + y8
+   # regressions
+   dem60 ~ ind60
+   dem65 ~ ind60 + dem60
+   # residual (co)variances
+   y1 ~~ y5
+   y2 ~~ y4 + y6
+   y3 ~~ y7
+   y4 ~~ y8
+   y6 ~~ y8
+ '
> fit <- sem(model, data=PoliticalDemocracy)
> coef(fit)

```

ind60=~x2	ind60=~x3	dem60=~y2	dem60=~y3	dem60=~y4	dem65=~y6
2.18036714	1.81851074	1.25674943	1.05771682	1.26479406	1.18569755
dem65=~y7	dem65=~y8	x1~~x1	x2~~x2	x3~~x3	y1~~y1
1.27951120	1.26594723	0.08154936	0.11980662	0.46670208	1.89140112
y5~~y1	y2~~y2	y4~~y2	y6~~y2	y3~~y3	y7~~y3
0.62366956	7.37297430	1.31303762	2.15291322	5.06753538	0.79500700
y4~~y4	y8~~y4	y5~~y5	y6~~y6	y8~~y6	y7~~y7
3.14779771	0.34822664	2.35096513	4.95394085	1.35617088	3.43142160
y8~~y8	dem60~ind60	dem65~ind60	dem65~dem60	ind60~~ind60	dem60~~dem60
3.25413086	1.48300197	0.57233362	0.83734605	0.44843768	3.95598568
dem65~~dem65					
0.17248060					

The `coef` function extracts the estimated values of the free parameters in the model, together with their names. Each name consists of three parts and reflects the part of the formula where the parameter was involved. The first part is the variable name that appears on the left-hand side of the formula. The middle part is the operator type of the formula, and the third part is the variable in the right-hand side of the formula that corresponds with the parameter.

If you want, you can provide custom parameter names by using the `label()` modifier. An example will make this clear:

```

> model <- '
+   # latent variable definitions
+   ind60 =~ x1 + x2 + label("myLabel")*x3
+   dem60 =~ y1 + y2 + y3 + y4
+   dem65 =~ y5 + y6 + y7 + y8
+   # regressions
+   dem60 ~ ind60
+   dem65 ~ ind60 + dem60
+   # residual (co)variances
+   y1 ~~ y5
+   y2 ~~ y4 + y6
+   y3 ~~ y7
+   y4 ~~ y8
+   y6 ~~ y8
+ '

```

The default name of the parameter associated with the factor loading of the `x3` indicator is by default `"ind60= x3"`. But the `label()` modifier will change it to the custom name `"myLabel"`.

5.4 Equality constraints

In some applications, it is useful to impose equality constraints on one or more otherwise free parameters. Consider again the three-factor H&S CFA model. Suppose a user has a priori reasons to believe that the factor loadings of the `x2` and `x3` indicators are equal to each other. Instead of estimating two free parameters, `lavaan` should only estimate a single free parameter, and use that value for both factor loadings. Another way of thinking about this is that the factor loading for the `x2` variable will be freely estimated, but that the factor

loading of the `x3` variable will be set equal to the factor loading of the `x2` variable. We call the factor loading for `x2` the ‘target parameter’, and the factor loading for `x3` the ‘constrained’ parameter. In the `lavaan` model syntax, we again need to use the pre-multiplication mechanism using a special function called `equal()`. The single argument of this function is the name of the target parameter. This is illustrated in the following syntax:

lavaan syntax

```
visual  =~ x1 + x2 + equal("visual=~x2")*x3
textual =~ x4 + x5 + x6
speed   =~ x7 + x8 + x9
```

The parameter corresponding to the factor loading of the `x2` variable is (automatically) called `"visual=~x2"`. By using the `equal()` modifier for `x3`, the corresponding parameter value will be set equal to the factor loading of `x2`. This mechanism can be used for any free parameter in a `lavaan` model.

6 Meanstructures and multiple groups

6.1 Bringing in the means

By and large, structural equation models are used to model the covariance matrix of the observed variables in a dataset. But in some applications, it is useful to bring in the means of the observed variables too. One way to do this is to explicitly refer to intercepts in the `lavaan` syntax. This can be done by including ‘intercept formulas’ in the model syntax. An intercept formula has the following form:

```
variable ~ 1
```

The left part of the expression contains the name of the observed or latent variable. The right part contains the number 1, representing the intercept. For example, in the three-factor H&S CFA model, we can add the intercepts of the observed variables as follows:

lavaan syntax

```
# three-factor model
visual  =~ x1 + x2 + x3
textual =~ x4 + x5 + x6
speed   =~ x7 + x8 + x9

# intercepts
x1 ~ 1
x2 ~ 1
x3 ~ 1
x4 ~ 1
x5 ~ 1
x6 ~ 1
x7 ~ 1
x8 ~ 1
x9 ~ 1
```

However, it is more convenient to omit the intercept formulas in the model syntax (unless you want to fix their values), and to add the `meanstructure = TRUE` argument in the `cfa` and `sem` function calls. For example, we can refit the three-factor H&S CFA model as follows:

```
> fit <- cfa(HS.model, data = HolzingerSwineford1939, meanstructure = TRUE)
> summary(fit)
```

Model converged normally after 35 iterations using ML

Minimum Function Chi-square	85.306
Degrees of freedom	24
P-value	0.0000

Estimate	Std.err	Z-value	P(> z)
----------	---------	---------	---------

```

Latent variables:
  visual =~
    x1          1.000
    x2          0.554    0.100    5.554    0.000
    x3          0.729    0.109    6.685    0.000
  textual =~
    x4          1.000
    x5          1.113    0.065   17.014    0.000
    x6          0.926    0.055   16.703    0.000
  speed =~
    x7          1.000
    x8          1.180    0.165    7.152    0.000
    x9          1.082    0.151    7.155    0.000

Latent covariances:
  visual ~~
    textual      0.408    0.074    5.552    0.000
    speed        0.262    0.056    4.660    0.000
  textual ~~
    speed        0.173    0.049    3.518    0.000

Latent means/intercepts:
  visual      0.000
  textual     0.000
  speed       0.000

Intercepts:
  x1          4.936    0.067   73.473    0.000
  x2          6.088    0.068   89.855    0.000
  x3          2.250    0.065   34.579    0.000
  x4          3.061    0.067   45.694    0.000
  x5          4.341    0.074   58.452    0.000
  x6          2.186    0.063   34.667    0.000
  x7          4.186    0.063   66.766    0.000
  x8          5.527    0.058   94.854    0.000
  x9          5.374    0.058   92.546    0.000

Latent variances:
  visual      0.809    0.145    5.564    0.000
  textual     0.979    0.112    8.737    0.000
  speed       0.384    0.086    4.451    0.000

Residual variances:
  x1          0.549    0.114    4.833    0.000
  x2          1.134    0.102   11.146    0.000
  x3          0.844    0.091    9.317    0.000
  x4          0.371    0.048    7.778    0.000
  x5          0.446    0.058    7.642    0.000
  x6          0.356    0.043    8.277    0.000
  x7          0.799    0.081    9.823    0.000
  x8          0.488    0.074    6.573    0.000
  x9          0.566    0.071    8.003    0.000

```

As you can see in the output, the model includes intercept parameters for both the observed and latent variables. By default, the latent variable intercepts (which in this case correspond to the latent *means*) are fixed to zero. Otherwise, the model would not be estimable. Note that the chi-square statistic and the number of degrees of freedom is the same as in the original (non-meanstructure) model. The reason is that we brought in some new data (a mean value for each of the 9 observed variables), but we also added 9 additional parameters to the model (an intercept for each of the 9 observed variables). The end result is an identical fit. In practice, the only reason why a user would add intercept-formulas in the model syntax, is because some constraints must be specified on them. For example, suppose that we wish to fix the intercepts of the variables `x1`, `x2`, `x3` and `x4` to, say, 0.5. We would write the model syntax as follows:

```
# three-factor model
visual  =~ x1 + x2 + x3
textual =~ x4 + x5 + x6
speed   =~ x7 + x8 + x9

# intercepts with fixed values
x1 ~ 0.5*1
x2 ~ 0.5*1
x3 ~ 0.5*1
x4 ~ 0.5*1
```

6.2 Multiple groups

The `lavaan` package has full support for multiple groups. To request a multiple group analysis, you need to add the name of the group variable in your dataset to the `group` argument in the `cfa` and `sem` function calls. By default, the same model is fitted in all groups. In the following example, we fit the H&S CFA model for the two schools (Pasteur and Grant-White).

```
> HS.model <- ' visual  =~ x1 + x2 + x3
+               textual =~ x4 + x5 + x6
+               speed   =~ x7 + x8 + x9 '
> fit <- cfa(HS.model, data=HolzingerSwineford1939, group="school")
> summary(fit)
```

Model converged normally after 56 iterations using ML

Minimum Function Chi-square	115.851
Degrees of freedom	48
P-value	0.0000

Chi-square for each group:

Grant-White	51.542
Pasteur	64.309

Group 1 [Grant-White]:

	Estimate	Std.err	Z-value	P(> z)
Latent variables:				
visual =~				
x1	1.000			
x2	0.736	0.155	4.760	0.000
x3	0.925	0.166	5.583	0.000
textual =~				
x4	1.000			
x5	0.990	0.087	11.418	0.000
x6	0.963	0.085	11.377	0.000
speed =~				
x7	1.000			
x8	1.226	0.187	6.569	0.000
x9	1.058	0.165	6.429	0.000
Latent covariances:				
visual ~~				
textual	0.408	0.098	4.153	0.000
speed	0.276	0.076	3.639	0.000
textual ~~				
speed	0.222	0.073	3.022	0.003
Latent variances:				
visual	0.604	0.160	3.762	0.000
textual	0.942	0.152	6.177	0.000

speed	0.461	0.118	3.910	0.000
-------	-------	-------	-------	-------

Residual variances:

x1	0.715	0.126	5.676	0.000
x2	0.899	0.123	7.339	0.000
x3	0.557	0.103	5.409	0.000
x4	0.315	0.065	4.870	0.000
x5	0.419	0.072	5.812	0.000
x6	0.406	0.069	5.880	0.000
x7	0.600	0.091	6.584	0.000
x8	0.401	0.094	4.248	0.000
x9	0.535	0.089	6.010	0.000

Group 2 [Pasteur]:

	Estimate	Std.err	Z-value	P(> z)
--	----------	---------	---------	---------

Latent variables:

visual =~

x1	1.000			
x2	0.394	0.122	3.220	0.001
x3	0.570	0.140	4.076	0.000

textual =~

x4	1.000			
x5	1.183	0.102	11.613	0.000
x6	0.875	0.077	11.421	0.000

speed =~

x7	1.000			
x8	1.125	0.277	4.058	0.000
x9	0.922	0.225	4.104	0.000

Latent covariances:

visual ~~

textual	0.479	0.106	4.531	0.000
speed	0.185	0.077	2.397	0.017

textual ~~

speed	0.182	0.069	2.628	0.009
-------	-------	-------	-------	-------

Latent variances:

visual	1.097	0.276	3.967	0.000
textual	0.894	0.150	5.963	0.000
speed	0.350	0.126	2.778	0.005

Residual variances:

x1	0.298	0.232	1.286	0.198
x2	1.334	0.158	8.464	0.000
x3	0.989	0.136	7.271	0.000
x4	0.425	0.069	6.138	0.000
x5	0.456	0.086	5.292	0.000
x6	0.290	0.050	5.780	0.000
x7	0.820	0.125	6.580	0.000
x8	0.510	0.116	4.406	0.000
x9	0.680	0.104	6.516	0.000

If you want to fix parameters, or provide starting values, you can use the same pre-multiplication techniques, but the single argument is now replaced by a vector of arguments, one for each group. For example:

lavaan syntax

```
HS.model <- ' visual =~          x1 +
              x2 +
              c(0.6, 0.8)*x3

              textual =~          x4 +
              start(c(1.2, 0.6))*x5 +
              x6'
```

```

speed    =~
                                x7 +
                                x8 +
                                label(c("x9.group1",
                                           "x9.group2"))*x9 '

```

In the definition of the latent factor `visual`, we have fixed the factor loading of the `x3` indicator to the value ‘0.6’ in the first group, and to the value ‘0.8’ in the second group. In the definition of the `textual` factor, two different starting values are provided for the `x5` indicator; one for each group. Finally, in the definition of the `speed` factor, we changed the labels of the parameters associated with the factor loading of the `x9` indicator. For the last modification, we can see the effect by requesting the values of the estimated parameters:

```

> fit <- cfa(HS.model, data = HolzingerSwineford1939, group = "school")
> coef(fit)

```

Grant-White.visual=~x2	Grant-White.textual=~x5
0.5880695	0.9904237
Grant-White.textual=~x6	Grant-White.speed=~x8
0.9617101	1.2276221
x9.group1	Grant-White.x1~~x1
1.0892346	0.5712898
Grant-White.x2~~x2	Grant-White.x3~~x3
0.9408347	0.6825272
Grant-White.x4~~x4	Grant-White.x5~~x5
0.3146871	0.4171026
Grant-White.x6~~x6	Grant-White.x7~~x7
0.4084114	0.6115976
Grant-White.x8~~x8	Grant-White.x9~~x9
0.4158980	0.5169119
Grant-White.visual~~visual	Grant-White.textual~~visual
0.8256916	0.4706483
Grant-White.speed~~visual	Grant-White.textual~~textual
0.3353609	0.9425314
Grant-White.speed~~textual	Grant-White.speed~~speed
0.2233279	0.4502038
Pasteur.visual=~x2	Pasteur.textual=~x5
0.5287217	1.1895757
Pasteur.textual=~x6	Pasteur.speed=~x8
0.8777832	1.1373990
x9.group2	Pasteur.x1~~x1
0.9525108	0.5387143
Pasteur.x2~~x2	Pasteur.x3~~x3
1.2742905	0.8786946
Pasteur.x4~~x4	Pasteur.x5~~x5
0.4304867	0.4496137
Pasteur.x6~~x6	Pasteur.x7~~x7
0.2893181	0.8308872
Pasteur.x8~~x8	Pasteur.x9~~x9
0.5134012	0.6699607
Pasteur.visual~~visual	Pasteur.textual~~visual
0.8208162	0.4023822
Pasteur.speed~~visual	Pasteur.textual~~textual
0.1742155	0.8892006
Pasteur.speed~~textual	Pasteur.speed~~speed
0.1784486	0.3390281

If multiple groups are involved, the ‘default’ parameter names include the name of the group. The labels for the `x9` indicator are changed to the custom names provided via the `label()` modifier.

6.2.1 Constraining a single parameter to be equal across groups

If you want to constrain one or more parameters to be equal across groups, we can again use the `equal()` modifier. For example, to constrain the factor loading of the `x3` indicator to be equal across groups, we can use the `equal()` modifier as follows:


```

> HS.model <- ' visual  =~ x1 + x2 +
+                  equal(c("", "Grant-White.visual=~x2")) *x3
+                  textual =~ x4 + x5 + x6
+                  speed   =~ x7 + x8 + x9 '

```

The first element of the `equal` modifier is the empty string: we do *not* want to impose any equality constraints on the factor loading in the first group. All we want is to set the value of the factor loading (of `x3`) in the second group (the Pasteur school) equal to the freely estimated value in first group (the Grant-White school). That is why we take `"Grant-White.visual= x2"` as the label of the ‘target’ parameter in this group.

6.2.2 Constraining groups of parameters to be equal across groups

Although the `equal()` modifier is very flexible, there is a more convenient way to impose equality constraints on a whole set of parameters (for example: all factor loadings, or all intercepts). We call these type of constraints *group constraints* and they can be specified by the `group.constraints` argument in the `cfa` or `sem` function call. For example, to constrain (all) the factor loadings to be equal across groups, you can proceed as follows:

```

> HS.model <- ' visual  =~ x1 + x2 + x3
+                  textual =~ x4 + x5 + x6
+                  speed   =~ x7 + x8 + x9 '
> fit <- cfa(HS.model, data=HolzingerSwineford1939, group="school",
+            group.constraints=c("loadings"))
> summary(fit)

```

Model converged normally after 42 iterations using ML

Minimum Function Chi-square	124.044
Degrees of freedom	54
P-value	0.0000

Chi-square for each group:

Grant-White	55.219
Pasteur	68.825

Group 1 [Grant-White]:

	Estimate	Std.err	Z-value	P(> z)
Latent variables:				
visual =~				
x1	1.000			
x2	0.599	0.100	5.979	0.000
x3	0.784	0.108	7.267	0.000
textual =~				
x4	1.000			
x5	1.083	0.067	16.049	0.000
x6	0.912	0.058	15.785	0.000
speed =~				
x7	1.000			
x8	1.201	0.155	7.738	0.000
x9	1.038	0.136	7.629	0.000

Latent covariances:				
visual ~~				
textual	0.437	0.099	4.423	0.000
speed	0.314	0.079	3.958	0.000
textual ~~				
speed	0.226	0.072	3.144	0.002

Latent variances:				
visual	0.722	0.161	4.490	0.000
textual	0.906	0.136	6.646	0.000
speed	0.475	0.109	4.347	0.000

Residual variances:

x1	0.645	0.127	5.084	0.000
x2	0.933	0.121	7.732	0.000
x3	0.605	0.096	6.282	0.000
x4	0.329	0.062	5.279	0.000
x5	0.384	0.073	5.270	0.000
x6	0.437	0.067	6.576	0.000
x7	0.599	0.090	6.651	0.000
x8	0.406	0.089	4.541	0.000
x9	0.532	0.086	6.202	0.000

Group 2 [Pasteur]:

	Estimate	Std.err	Z-value	P(> z)
Latent variables:				
visual =~				
x1	1.000			
x2	0.599			
x3	0.784			
textual =~				
x4	1.000			
x5	1.083			
x6	0.912			
speed =~				
x7	1.000			
x8	1.201			
x9	1.038			
Latent covariances:				
visual ~~				
textual	0.416	0.097	4.271	0.000
speed	0.169	0.064	2.643	0.008
textual ~~				
speed	0.176	0.061	2.882	0.004
Latent variances:				
visual	0.805	0.171	4.714	0.000
textual	0.913	0.137	6.651	0.000
speed	0.305	0.078	3.920	0.000
Residual variances:				
x1	0.551	0.137	4.010	0.000
x2	1.258	0.155	8.117	0.000
x3	0.882	0.128	6.884	0.000
x4	0.434	0.070	6.238	0.000
x5	0.508	0.082	6.229	0.000
x6	0.266	0.050	5.294	0.000
x7	0.849	0.114	7.468	0.000
x8	0.515	0.095	5.409	0.000
x9	0.658	0.096	6.865	0.000

More ‘group’ constraints can be added. In addition to the factor loadings, you can also constrain the "intercepts" of the observed variables, "means" of latent variables, and "residuals" (residual variances of observed variables) to be equal across groups, simply by adding them to the `group.constraints` argument. If you omit the `group.constraints` arguments, all parameters are freely estimated in each group (but the model structure is the same).

6.2.3 Measurement Invariance

If you are interested in testing the measurement invariance of a CFA model across several groups, you can use the `measurement.invariance` function which performs a number of multiple group analyses in a particular sequence, with increasingly more restrictions on the parameters. Each model is compared to the baseline model and the previous model using chi-square difference tests. In addition, the difference in the `cfi` fit measure is also shown. Although the current implementation of the function is still a bit primitive, it does illustrate

how the various components of the `lavaan` package can be used as building blocks for constructing higher level functions (such as the `measurement.invariance` function), something that is often very hard to accomplish with commercial software.

```
> measurement.invariance(HS.model, data = HolzingerSwineford1939,
+   group = "school")
```

Measurement invariance tests:

Model 1: configural invariance

chisq	df	pvalue	cfi	tli	rmsea	bic
115.851	48.000	0.000	0.924	0.900	0.097	7604.094

Model 2: weak invariance (equal loadings):

chisq	df	pvalue	cfi	tli	rmsea	bic
124.044	54.000	0.000	0.922	0.909	0.093	7578.043

[Model 1 versus model 2]

delta.chisq	df	p.value	delta.cfi
8.19	6	0.22436	0.0025

Model 3: strong invariance (equal loadings + equal intercepts):

chisq	df	pvalue	cfi	tli	rmsea	bic
164.103	60.000	0.000	0.884	0.878	0.107	7686.588

[Model 1 versus model 3]

delta.chisq	df	p.value	delta.cfi
48.25	12	0.00000	0.0405

[Model 2 versus model 3]

delta.chisq	df	p.value	delta.cfi
40.06	6	0.00000	0.0381

Model 4: equal loadings + intercepts + means:

chisq	df	pvalue	cfi	tli	rmsea	bic
204.605	63.000	0.000	0.855	0.835	0.122	7709.969

[Model 1 versus model 4]

delta.chisq	df	p.value	delta.cfi
88.75	15	0.00000	0.0688

[Model 3 versus model 4]

delta.chisq	df	p.value	delta.cfi
40.50	3	0.00000	0.0283

7 Growth curve models

Another important type of latent variable models are latent growth curve models. Growth modeling is often used to analyze longitudinal or developmental data. In this type of data, an outcome measure is measured on several occasions, and we want to study the change over time. In many cases, the trajectory over time can be modeled as a simple linear or quadratic curve. Random effects are used to capture individual differences. The random effects are conveniently represented by (continuous) latent variables, often called *growth factors*. In the example below, we use an artificial toy dataset called `Demo.growth` where a score (say, a standardized score on an reading ability scale) is measured on 4 time points. To fit a linear growth model for these four time points, we need to specify a model with two latent variables: a random intercept, and a random slope:

lavaan syntax

```
# linear growth model with 4 timepoints
# intercept and slope with fixed coefficients
i =~ 1*t1 + 1*t2 + 1*t3 + 1*t4
s =~ 0*t1 + 1*t2 + 2*t3 + 3*t4
```

In this model, we have fixed all the coefficients of the growth functions. To fit this model, the `lavaan` package provides a special `growth` function:

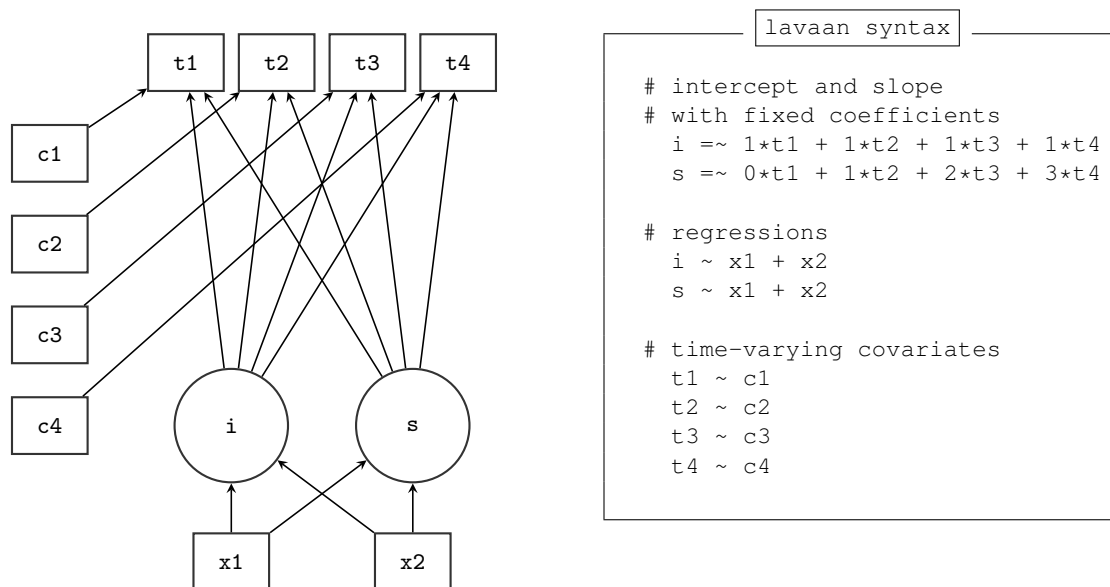
```
> model <- ' i =~ 1*t1 + 1*t2 + 1*t3 + 1*t4
+           s =~ 0*t1 + 1*t2 + 2*t3 + 3*t4 '
> fit <- growth(model, data=Demo.growth)
> summary(fit)
```

Model converged normally after 37 iterations using ML

Minimum Function Chi-square	8.069
Degrees of freedom	5
P-value	0.1525

	Estimate	Std.err	Z-value	P(> z)
Latent variables:				
i =~				
t1	1.000			
t2	1.000			
t3	1.000			
t4	1.000			
s =~				
t1	0.000			
t2	1.000			
t3	2.000			
t4	3.000			
Latent covariances:				
i ~~				
s	0.618	0.071	8.686	0.000
Latent means/intercepts:				
i	0.615	0.077	8.007	0.000
s	1.006	0.042	24.076	0.000
Intercepts:				
t1	0.000			
t2	0.000			
t3	0.000			
t4	0.000			
Latent variances:				
i	1.932	0.173	11.194	0.000
s	0.587	0.052	11.336	0.000
Residual variances:				
t1	0.595	0.086	6.944	0.000
t2	0.676	0.061	11.061	0.000
t3	0.635	0.072	8.761	0.000
t4	0.508	0.124	4.090	0.000

Technically, the `growth` function is almost identical to the `sem` function. But a meanstructure is automatically assumed, and the observed intercepts are fixed to zero by default, while the latent variable intercepts/means are freely estimated. A slightly more complex model adds two regressors (`x1` and `x2`) that influence the latent growth factors. In addition, a time-varying covariate that influences the outcome measure at the four time points has been added to the model. A graphical representation of this model together with the corresponding `lavaan` syntax is presented below.



For ease of copy/pasting, the complete R code needed to specify and fit this linear growth model with a time-varying covariate is printed again below:

R code

```

# a linear growth model with a time-varying covariate

model <- '
# intercept and slope with fixed coefficients
i =~ 1*t1 + 1*t2 + 1*t3 + 1*t4
s =~ 0*t1 + 1*t2 + 2*t3 + 3*t4

# regressions
i ~ x1 + x2
s ~ x1 + x2

# time-varying covariates
t1 ~ c1
t2 ~ c2
t3 ~ c3
t4 ~ c4
'

fit <- growth(model, data=Demo.growth)
summary(fit)

```

8 Additional information

8.1 Using a covariance matrix as input

If you have no full dataset, but you do have a sample covariance matrix, you can still fit your model. If you need a meanstructure, you will need to provide a mean vector too. Importantly, you also need to specify the number of observations that were used to compute the sample moments. The following example illustrates the use of a sample covariance matrix as input:

```

> wheaton.cov <- matrix(c(
+   11.834,    0,    0,    0,    0,    0,
+   6.947,    9.364,    0,    0,    0,    0,
+   6.819,    5.091,  12.532,    0,    0,    0,
+   4.783,    5.028,    7.495,    9.986,    0,    0,

```

```

+      -3.839, -3.889, -3.841, -3.625, 9.610, 0,
+      -21.899, -18.831, -21.748, -18.775, 35.522, 450.288),
+      6, 6, byrow=TRUE)
> colnames(wheaton.cov) <- rownames(wheaton.cov) <-
+   c("anomia67", "powerless67", "anomia71",
+     "powerless71", "education", "sei")
> wheaton.model <- '
+   # measurement model
+   ses      =~ education + sei
+   alien67 =~ anomia67 + powerless67
+   alien71 =~ anomia71 + powerless71
+
+   # equations
+   alien71 ~ alien67 + ses
+   alien67 ~ ses
+
+   # correlated residuals
+   anomia67 ~~ anomia71
+   powerless67 ~~ powerless71
+ '
> fit <- sem(wheaton.model, sample.cov=wheaton.cov, sample.nobs=932)
> summary(fit, standardized=TRUE)

```

Model converged normally after 112 iterations using ML

Minimum Function Chi-square	4.735
Degrees of freedom	4
P-value	0.3156

	Estimate	Std.err	Z-value	P(> z)	Std.lv	Std.all
Latent variables:						
ses =~						
education	1.000				2.607	0.842
sei	5.219	0.422	12.364	0.000	13.609	0.642
alien67 =~						
anomia67	1.000				2.663	0.774
powerless67	0.979	0.062	15.895	0.000	2.606	0.852
alien71 =~						
anomia71	1.000				2.850	0.805
powerless71	0.922	0.059	15.498	0.000	2.628	0.832
Regressions:						
alien67 ~						
ses	-0.575	0.056	-10.195	0.000	-0.563	-0.563
alien71 ~						
ses	-0.227	0.052	-4.334	0.000	-0.207	-0.207
alien67	0.607	0.051	11.898	0.000	0.567	0.567
Residual covariances:						
anomia67 ~~						
anomia71	1.623	0.314	5.176	0.000	1.623	0.133
powerless67 ~~						
powerless71	0.339	0.261	1.298	0.194	0.339	0.035
Latent variances:						
ses	6.798	0.649	10.475	0.000	1.000	1.000
Residual variances:						
education	2.801	0.507	5.525	0.000	2.801	0.292
sei	264.597	18.126	14.597	0.000	264.597	0.588
anomia67	4.731	0.453	10.441	0.000	4.731	0.400
powerless67	2.563	0.403	6.359	0.000	2.563	0.274
anomia71	4.399	0.515	8.542	0.000	4.399	0.351
powerless71	3.070	0.434	7.070	0.000	3.070	0.308
alien67	4.841	0.467	10.359	0.000	0.683	0.683

Only the lower half elements of the covariance matrix (including the diagonal) is used. The rownames (and optionally the colnames) must contain the names of the observed variables that are used in the model syntax. If you have multiple groups, the `sample.cov` argument must be a list containing the sample variance-covariance matrix of each group as a separate element in the list. If a meanstructure is needed, the `sample.mean` argument must be a list containing the sample means of each group. Finally, the `sample.nobs` argument can be either a list or a integer vector containing the number of observations for each group.

8.2 Estimators, Standard errors and Missing Values

8.2.1 Estimators

The default estimator in the `lavaan` package is maximum likelihood (`estimator = "ML"`). Alternative estimators currently available in `lavaan` are:

- "GLS" for generalized least squares
- "WLS" for weighted least squares (sometimes called ADF estimation)
- "MLM" for maximum likelihood estimation with robust standard errors and a Satorra-Bentler scaled test statistic.

If maximum likelihood estimation is used ("ML" or "MLM"), the default behavior of `lavaan` is to base the analysis on the so-called *biased* sample covariance matrix, where the elements are divided by n instead of $n - 1$. This is done internally, and should not be done by the user. In addition, the chi-square statistic is computed by multiplying the minimum function value with a factor n (instead of $n - 1$). This is similar to the Mplus program. If you prefer to use an unbiased covariance, and $n - 1$ as the multiplier to compute the chi-square statistic, you need to specify the `mimic.Mplus=FALSE` argument when calling the fitting functions.

8.2.2 Missing values

If the data contain missing values, the default behavior is listwise deletion. If the missing mechanism is MCAR (missing completely at random) or MAR (missing at random), the `lavaan` package provides case-wise (or 'full information') maximum likelihood estimation. You can 'turn' this feature on, by using the argument `na.rm=FALSE` when calling the fitting function. An unrestricted (h1) model will automatically be estimated, so that all common fit indices are available.

8.2.3 Standard Errors

Standard errors are (by default) based on the expected information matrix. The only exception is when data are missing and full information ML is used (via `na.rm=FALSE`). In this case, the observed information matrix is used to compute the standard errors. The user can change this behavior by using the `information` argument, which can be set to "expected" or "observed". If the "MLM" estimator is used, the standard errors are based on the expected information matrix and corrected using the Satorra-Bentler approach.

8.3 Modification Indices

Modification indices can be requested by adding the `modindices=TRUE` argument in the `summary` call. For example:

```
> fit <- cfa(HS.model, data = HolzingerSwineford1939)
> summary(fit, modindices = TRUE)
```

Model converged normally after 35 iterations using ML

Minimum Function Chi-square	85.306
Degrees of freedom	24
P-value	0.0000

	Estimate	Std.err	Z-value	P(> z)
Latent variables:				
visual =~				
x1	1.000			

x2	0.554	0.100	5.554	0.000
x3	0.729	0.109	6.685	0.000
textual =~				
x4	1.000			
x5	1.113	0.065	17.014	0.000
x6	0.926	0.055	16.703	0.000
speed =~				
x7	1.000			
x8	1.180	0.165	7.152	0.000
x9	1.082	0.151	7.155	0.000
Latent covariances:				
visual ~~				
textual	0.408	0.074	5.552	0.000
speed	0.262	0.056	4.660	0.000
textual ~~				
speed	0.173	0.049	3.518	0.000
Latent variances:				
visual	0.809	0.145	5.564	0.000
textual	0.979	0.112	8.737	0.000
speed	0.384	0.086	4.451	0.000
Residual variances:				
x1	0.549	0.114	4.833	0.000
x2	1.134	0.102	11.146	0.000
x3	0.844	0.091	9.317	0.000
x4	0.371	0.048	7.778	0.000
x5	0.446	0.058	7.642	0.000
x6	0.356	0.043	8.277	0.000
x7	0.799	0.081	9.823	0.000
x8	0.488	0.074	6.573	0.000
x9	0.566	0.071	8.003	0.000

Modification Indices:

Parameter label	M.I.	E.P.C.	Std.lv	Std.all
visual=~x4	1.211	0.077	0.069	0.059
visual=~x5	7.441	-0.210	-0.189	-0.147
visual=~x6	2.843	0.111	0.100	0.092
visual=~x7	18.631	-0.422	-0.380	-0.349
visual=~x8	4.295	-0.210	-0.189	-0.187
visual=~x9	36.411	0.577	0.519	0.515
textual=~x1	8.903	0.350	0.347	0.297
textual=~x2	0.017	-0.011	-0.011	-0.010
textual=~x3	9.151	-0.272	-0.269	-0.238
textual=~x7	0.098	-0.021	-0.021	-0.019
textual=~x8	3.359	-0.121	-0.120	-0.118
textual=~x9	4.796	0.138	0.137	0.136
speed=~x1	0.014	0.024	0.015	0.013
speed=~x2	1.580	-0.198	-0.123	-0.105
speed=~x3	0.716	0.136	0.084	0.075
speed=~x4	0.003	-0.005	-0.003	-0.003
speed=~x5	0.201	-0.044	-0.027	-0.021
speed=~x6	0.273	0.044	0.027	0.025
x2~~x1	3.606	-0.184	-0.184	-0.134
x3~~x1	0.935	-0.139	-0.139	-0.105
x3~~x2	8.532	0.218	0.218	0.164
x4~~x1	3.554	0.078	0.078	0.058
x4~~x2	0.534	-0.034	-0.034	-0.025
x4~~x3	0.142	-0.016	-0.016	-0.012
x5~~x1	0.522	-0.033	-0.033	-0.022
x5~~x2	0.023	-0.008	-0.008	-0.005

x5~~x3	7.858	-0.130	-0.130	-0.089
x5~~x4	2.534	0.186	0.186	0.124
x6~~x1	0.048	0.009	0.009	0.007
x6~~x2	0.785	0.039	0.039	0.031
x6~~x3	1.855	0.055	0.055	0.044
x6~~x4	6.221	-0.235	-0.235	-0.185
x6~~x5	0.916	0.101	0.101	0.072
x7~~x1	5.420	-0.129	-0.129	-0.102
x7~~x2	8.918	-0.183	-0.183	-0.143
x7~~x3	0.638	-0.044	-0.044	-0.036
x7~~x4	5.920	0.098	0.098	0.078
x7~~x5	1.233	-0.049	-0.049	-0.035
x7~~x6	0.259	-0.020	-0.020	-0.017
x8~~x1	0.634	-0.041	-0.041	-0.035
x8~~x2	0.054	-0.012	-0.012	-0.010
x8~~x3	0.059	-0.012	-0.012	-0.011
x8~~x4	3.805	-0.069	-0.069	-0.059
x8~~x5	0.347	0.023	0.023	0.018
x8~~x6	0.275	0.018	0.018	0.016
x8~~x7	34.145	0.536	0.536	0.488
x9~~x1	7.335	0.138	0.138	0.117
x9~~x2	1.895	0.075	0.075	0.063
x9~~x3	4.126	0.102	0.102	0.089
x9~~x4	0.196	-0.016	-0.016	-0.014
x9~~x5	0.999	0.040	0.040	0.031
x9~~x6	0.097	-0.011	-0.011	-0.010
x9~~x7	5.183	-0.187	-0.187	-0.170
x9~~x8	14.946	-0.423	-0.423	-0.415

Modification indices are printed out for each nonfree (or nonredundant) parameter. The modification indices are supplemented by the expected parameter change values (column E.P.C.). The last two columns are the standardized and completely standardized EPC values respectively.

8.4 Extracting information from a fitted model

If you want to peek inside a fitted `semModel` object (the object that is returned by a call to `cfa`, `sem` or `growth`), or you want to ‘extract’ specific information from a fitted object, you can use the `inspect` function, with a variety of options. By default, calling `inspect` on a fitted `semModel` object returns a list of the model matrices that are used internally to represent the model. The free parameters are nonzero integers.

```
> inspect(fit)

$lambda
      visual textual speed
x1         0         0     0
x2         1         0     0
x3         2         0     0
x4         0         0     0
x5         0         3     0
x6         0         4     0
x7         0         0     0
x8         0         0     5
x9         0         0     6

$theta
      x1 x2 x3 x4 x5 x6 x7 x8 x9
x1    7  0  0  0  0  0  0  0  0
x2    0  8  0  0  0  0  0  0  0
x3    0  0  9  0  0  0  0  0  0
x4    0  0  0 10  0  0  0  0  0
x5    0  0  0  0 11  0  0  0  0
x6    0  0  0  0  0 12  0  0  0
x7    0  0  0  0  0  0 13  0  0
x8    0  0  0  0  0  0  0 14  0
x9    0  0  0  0  0  0  0  0 15
```

```
$psi
      visual textual speed
visual      16      0      0
textual     17     19      0
speed       18     20     21
```

To see the starting values of parameters in each model matrix, type

```
> inspect(fit, what = "start")
```

```
$lambda
      visual textual speed
x1         1         0         0
x2         1         0         0
x3         1         0         0
x4         0         1         0
x5         0         1         0
x6         0         1         0
x7         0         0         1
x8         0         0         1
x9         0         0         1

$theta
      x1      x2      x3      x4      x5      x6      x7
x1 0.6814489 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
x2 0.0000000 0.6931949 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
x3 0.0000000 0.0000000 0.6395572 0.0000000 0.0000000 0.0000000 0.0000000
x4 0.0000000 0.0000000 0.0000000 0.6775834 0.0000000 0.0000000 0.0000000
x5 0.0000000 0.0000000 0.0000000 0.0000000 0.8326592 0.0000000 0.0000000
x6 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.6001731 0.0000000
x7 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.5935416
x8 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
x9 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
      x8      x9
x1 0.0000000 0.0000000
x2 0.0000000 0.0000000
x3 0.0000000 0.0000000
x4 0.0000000 0.0000000
x5 0.0000000 0.0000000
x6 0.0000000 0.0000000
x7 0.0000000 0.0000000
x8 0.5126947 0.0000000
x9 0.0000000 0.5091936
```

```
$psi
      visual textual speed
visual     0.05     0.00  0.00
textual     0.00     0.05  0.00
speed       0.00     0.00  0.05
```

To extract a single fit index, say, CFI from a fitted model, you can use

```
> fit.indices <- inspect(fit, what = "fit")
> fit.indices["cfi"]
```

```
      cfi
0.9305597
```

The `inspect` function always returns the information as a vector or a list, so that the information can be captured for further processing. For more inspect options, see the help page for the `semModel` class which you can find by typing the following:

```
> class?semModel
```

Other ‘extractor’ functions are `coef`, `fitted.values`, `residuals` and its alias `resid`. We have already seen that `coef` returns the values of the free parameters (as a vector). The `fitted.values` function returns a list containing the *implied* covariance matrix and mean vector. The `residuals` and `resid` functions return a list containing the raw difference between the observed and implied covariance matrix and mean vector.

9 Known Issues

There are a number of known issues with the current beta version of the `lavaan` package. If you can help us out with any of these, we would be very grateful.

~= lavaan version 0.3-1 (march 2010) ~=

New issues:

- * multigroup WLS in Mplus: if each group has its own set of free parameters, the X2 of the first group is not same, as if you would fit the model separately using that group alone; strangely, the X2 of the second group is ok... Is this a bug in Mplus? (confirmed in 4.1 and 5.21)
- * In Mplus: SRMR index is (sometimes much) smaller if 'information=observed' option is used? (when?)

Old Issues:

- * Satorra-Bentler correction if `mimic.Mplus=TRUE`:
 - the standard errors and scaled chi-square test statistic are not the same as in Mplus (4.1 and 5.21); Mplus must be doing something different here, but what?Note: if `mimic.Mplus=FALSE`, the results are the same as in EQS.
- * Modification indices:
 - EPC's for equality constraints are not the same as in Mplus (for example, the NU EPC's in Mplus example 5.9); I need a reference containing the proper formulas!
 - the MI's for multiple group models with equality constraints are not identical to Mplus (but mostly close); again, a reference containing some formulas would be useful
- * SRMR value is slightly off in multiple-groups analyses with equality constraints (for example in ex5.14 with `equality.constraints=c("loadings", "intercepts", "means", "residuals")`, the SRMR in `semplus` is 0.133, while Mplus reports 0.142)

10 New features and changes compared to 'semplus 0.9-10' (december 2009)

~= lavaan version 0.3-1 (march 2010) ~=

User-visible changes (compared to `semplus 0.9-1 december 2009` version):

- * the name of the package has changed to `'lavaan'` (for latent variable analysis)
- * the `'ML.N'` option is replaced by a `'mimic.Mplus'` option; if `TRUE`, an attempt is made to mimic Mplus results as much as possible; if `FALSE`, results are more similar to EQS/LISREL results
- * if `do.fit=FALSE`, a full summary (including standard errors) is now available
- * if a correlation matrix is supplied (instead of a covariance matrix), only a (big) warning is now spit out (instead of an error and stopping)

New features (compared to `semplus 0.9-1 december 2009` version):

- * the model syntax can now be specified as a string literal enclosed in single quotes, allowing for arbitrary blank lines and comments; this will be the preferred way to provide a model syntax; the `specfiy.Model()` function will be deprecated in the next release

- * multiple values are now accepted within pre-multiplication commands when analyzing multiple groups; for example: `"start(c(0.5, 0.8))"` will give different starting values for each of the two groups
- * in a multiple group analysis, the sample moments can be provided using a list
- * 'automatic' naming of free parameters is now group-dependent
- * using `NA*x` in a formula forces the corresponding parameter to be free
- * a new modifier 'label' can now be used to specify custom labels for parameters, eg. `f =~ x1 + x2 + label("mylabel")*x3`
- * added 'information' argument to the `cfa/sem/growth` functions, so that the user can choose between the 'expected' or the 'observed' information matrix to be used when calculating standard errors; the observed information matrix is currently computed using a numerical approximation (not analytically); this produces accurate results, but is fairly slow
- * if `na.rm=FALSE` and `estimator="ML"`, full information ML (FIML) is used to estimate the parameters; a new 'missing' slot is provided in the `Sample` slot of a fitted object, containing information about the missing patterns and their sufficient statistics

Bug fixes (compared to `semplus` 0.9-1 december 2009 version):

- * the `std.lv=TRUE` argument is now working again
- * fixed two bugs in the `specify.Model()` function:
 - the comment character '#' can now appear in the first column
 - avoid confusion with `lv` defs containing `equal("f1=~x1")` statements
 but the `specify.Model()` function is deprecated and will be removed in the next release
- * fixed WLS + meanstructure issue: function value of 1st iteration is now identical to `Mplus` 4.1
- * symmetric matrices returned by `inspect()` are fully symmetric (not only showing the lower half)
- * Satorra-Bentler correction: if `mimic.Mplus=FALSE`, the scaled chi-square statistic and standard errors should now be identical to `EQS` in all cases
- * corrected the nomenclature of `H0` and `H1` in the output of `'summary(fit, fit.measures=TRUE)'`
- * Residual covariances in 'summary' output does not show the user-fixed residual variances anymore

11 Report a bug, or give use feedback

If you have found a bug, or something unpleasant happened, please let us now. You can send an email to Yves.Rosseel@UGent.be. Start the subject line with `[lavaan]`, and it will get the proper attention. To help us with your problem (and fix our bugs), we need two types of information from you:

1. a detailed description of the problem: what happened, which error message or warning message did you see, and when does it occur. If possible at all, provide a reproducible example of the syntax that generated the error.
2. the output of the following command in `R`:

```
> sessionInfo()
```

This will show vital information about your platform, the version of R, and other stuff that might help us with identifying the problem.

We also welcome all suggestions, both on the software and the documentation.

A Examples from the Mplus User's Guide

Below, we provide some examples of *lavaan* model syntax to mimic the examples in the Mplus User's guide. The datafiles can be downloaded from <http://www.statmodel.com/ugexcerpts.shtml>.

A.1 Chapter 3: Regression and Path Analysis

```
# ex3.1
Data <- read.table("ex3.1.dat")
names(Data) <- c("y1", "x1", "x2")

model.ex3.1 <- ' y1 ~ x1 + x2 '
fit <- sem(model.ex3.1, data=Data)
summary(fit, standardized=TRUE, fit.measures=TRUE)

# ex3.11
Data <- read.table("ex3.11.dat")
names(Data) <- c("y1", "y2", "y3",
                "x1", "x2", "x3")

model.ex3.11 <- ' y1 ~ x1 + x2 + x3
                 y2 ~ x1 + x2 + x3
                 y3 ~ y1 + y2 + x2 '

fit <- sem(model.ex3.11, data=Data)
summary(fit, standardized=TRUE, fit.measures=TRUE)
```

A.2 Chapter 5: Confirmatory factor analysis and structural equation modeling

```
# ex5.1
Data <- read.table("ex5.1.dat")
names(Data) <- paste("y", 1:6, sep="")

model.ex5.1 <- ' f1 =~ y1 + y2 + y3
                 f2 =~ y4 + y5 + y6 '

fit <- cfa(model.ex5.1, data=Data)
summary(fit, standardized=TRUE, fit.measures=TRUE)

# ex5.6
Data <- read.table("ex5.6.dat")
names(Data) <- paste("y", 1:12, sep="")

model.ex5.6 <- ' f1 =~ y1 + y2 + y3
                 f2 =~ y4 + y5 + y6
                 f3 =~ y7 + y8 + y9
                 f4 =~ y10 + y11 + y12
                 f5 =~ f1 + f2 + f3 + f4 '

fit <- cfa(model.ex5.6, data=Data, estimator="ML")
summary(fit, standardized=TRUE, fit.measures=TRUE)

# ex5.8
Data <- read.table("ex5.8.dat")
```

```

names(Data) <- c(paste("y", 1:6, sep=""), paste("x", 1:3, sep=""))

model.ex5.8 <- ' f1 =~ y1 + y2 + y3
                f2 =~ y4 + y5 + y6
                f1 ~ x1 + x2 + x3
                f2 ~ x1 + x2 + x3 '

fit <- cfa(model.ex5.8, data=Data, estimator="ML")
summary(fit, standardized=TRUE, fit.measures=TRUE)

# ex5.9
Data <- read.table("ex5.9.dat")
names(Data) <- c("y1a", "y1b", "y1c", "y2a", "y2b", "y2c")

model.ex5.9 <- ' f1 =~ 1*y1a + 1*y1b + 1*y1c
                f2 =~ 1*y2a + 1*y2b + 1*y2c
                y1a ~ 1
                y1b ~ equal("y1a~1") * 1
                y1c ~ equal("y1a~1") * 1
                y2a ~ 1
                y2b ~ equal("y2a~1") * 1
                y2c ~ equal("y2a~1") * 1 '

fit <- cfa(model.ex5.9, data=Data)
summary(fit, standardized=TRUE, fit.measures=TRUE)

# ex5.11
Data <- read.table("ex5.11.dat")
names(Data) <- paste("y", 1:12, sep="")

model.ex5.11 <- ' f1 =~ y1 + y2 + y3
                f2 =~ y4 + y5 + y6
                f3 =~ y7 + y8 + y9
                f4 =~ y10 + y11 + y12
                f3 ~ f1 + f2
                f4 ~ f3 '

fit <- sem(model.ex5.11, data=Data, estimator="ML")
summary(fit, standardized=TRUE, fit.measures=TRUE)

# ex5.14
Data <- read.table("ex5.14.dat")
names(Data) <- c("y1", "y2", "y3", "y4", "y5", "y6", "x1", "x2", "x3", "g")

model.ex5.14 <- ' f1 =~ y1 + equal(c("", "1.f1=~y2"))*y2 + y3
                f2 =~ y4 + equal(c("", "1.f2=~y5"))*y5
                   + equal(c("", "1.f2=~y6"))*y6
                f1 ~ x1 + x2 + x3
                f2 ~ x1 + x2 + x3 '

fit <- cfa(model.ex5.14, data=Data, group="g", meanstructure=FALSE)
summary(fit, standardized=TRUE, fit.measures=TRUE)

# ex5.15
Data <- read.table("ex5.15.dat")
names(Data) <- c("y1", "y2", "y3", "y4", "y5", "y6", "x1", "x2", "x3", "g")

model.ex5.14 <- ' f1 =~ y1 + equal(c("", "1.f1=~y2"))*y2 + y3
                f2 =~ y4 + equal(c("", "1.f2=~y5"))*y5
                   + equal(c("", "1.f2=~y6"))*y6
                f1 ~ x1 + x2 + x3
                f2 ~ x1 + x2 + x3

```

```

f1 ~ c(0,NA)*1
f2 ~ c(0,NA)*1
y1 ~ equal(c("", "1.y1~1"))*1
y2 ~ equal(c("", "1.y2~1"))*1
y3 ~ 1
y4 ~ equal(c("", "1.y4~1"))*1
y5 ~ equal(c("", "1.y5~1"))*1
y6 ~ equal(c("", "1.y6~1"))*1 '

```

```

fit <- cfa(model.ex5.14, data=Data, group="g", meanstructure=TRUE)
summary(fit, standardized=TRUE, fit.measures=TRUE)

```

A.3 Chapter 6: Growth modeling

```

# 6.1
Data <- read.table("ex6.1.dat")
names(Data) <- c("y11", "y12", "y13", "y14")

model.ex6.1 <- ' i =~ 1*y11 + 1*y12 + 1*y13 + 1*y14
                s =~ 0*y11 + 1*y12 + 2*y13 + 3*y14 '

```

```

fit <- growth(model.ex6.1, data=Data)
summary(fit, standardized=TRUE, fit.measures=TRUE)

```

```

#6.8
Data <- read.table("ex6.8.dat")
names(Data) <- c("y11", "y12", "y13", "y14")

model.ex6.8 <- '
  i =~ 1*y11 + 1*y12 + 1*y13 + 1*y14
  s =~ 0*y11 + 1*y12 + start(2)*y13 + start(3)*y14
,

fit <- growth(model.ex6.8, data=Data)
summary(fit, standardized=TRUE, fit.measures=TRUE)

```

```

#6.9
Data <- read.table("ex6.9.dat")
names(Data) <- c("y11", "y12", "y13", "y14")

model.ex6.9 <- '
  i =~ 1*y11 + 1*y12 + 1*y13 + 1*y14
  s =~ 0*y11 + 1*y12 + 2*y13 + 3*y14
  q =~ 0*y11 + 1*y12 + 4*y13 + 9*y14
,

```

```

fit <- growth(model.ex6.9, data=Data)
summary(fit, standardized=TRUE, fit.measures=TRUE)

```

```

#6.10
Data <- read.table("ex6.10.dat")
names(Data) <- c("y11", "y12", "y13", "y14", "x1", "x2", "a31", "a32", "a33", "a34")

model.ex6.10 <- '
  i =~ 1*y11 + 1*y12 + 1*y13 + 1*y14
  s =~ 0*y11 + 1*y12 + 2*y13 + 3*y14
  i ~ x1 + x2
  s ~ x1 + x2
  y11 ~ a31
  y12 ~ a32
  y13 ~ a33

```

```

    y14 ~ a34
  ,

fit <- growth(model.ex6.10, data=Data)
summary(fit, standardized=TRUE, fit.measures=TRUE)

#6.11
Data <- read.table("ex6.11.dat")
names(Data) <- c("y1", "y2", "y3", "y4", "y5")

modelex6.11 <- '
  i   =~ 1*y1 + 1*y2 + 1*y3 + 1*y4 + 1*y5
  s1  =~ 0*y1 + 1*y2 + 2*y3 + 2*y4 + 2*y5
  s2  =~ 0*y1 + 0*y2 + 0*y3 + 1*y4 + 2*y5
  ,

fit <- growth(modelex6.11, data=Data)
summary(fit, standardized=TRUE, fit.measures=TRUE)

```